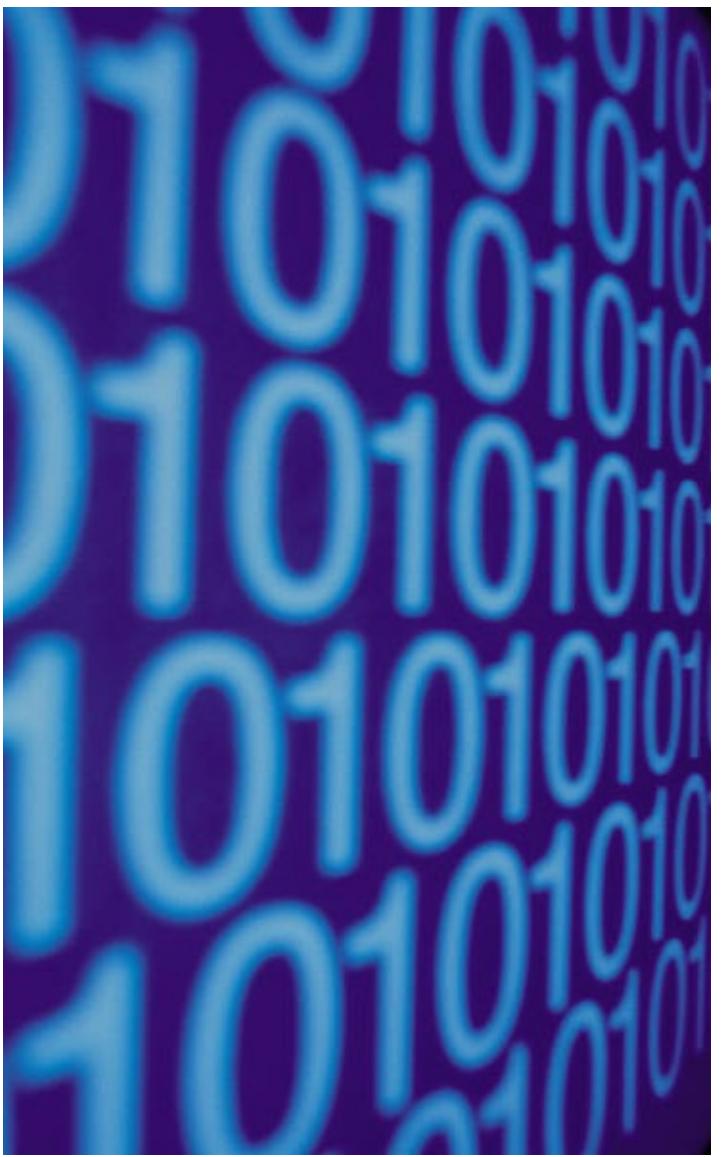


CPU, Memory and Processes

ECE568 – Lecture 2
Courtney Gibson, P.Eng.
University of Toronto ECE



Memory

Data Representation

Data Representation

As developers, our primary job is to manipulate numbers.



```

size_t
void *
int
void *
pthread_t
pthread_attr_t
int
struct main_args
int *

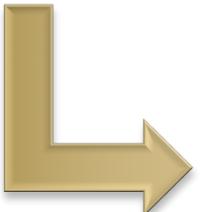
stackSize = STACK_SIZE;
targetStack = (void *)STACK_LOCATION;
pagesize = getpagesize();
mmap_result = NULL;
lab_main_thread_id;
pthread_attr;
rc = 0;
args = { argc , argv };
lab_main_return = NULL;

// Check that our target stack size is an integer multiple of pagesize
stacksize += pagesize - ( stackSize % pagesize );

// Allocate a new stack at a fixed location
mmap_result = mmap(targetStack, ( stackSize + pageSize ),
PROT_READ|PROT_WRITE|PROT_EXEC, MAP_ANON|MAP_FIXED);
assert ( mmap_result != MAP_FAILED );

// Create a new pthread to run lab_main()
rc = pthread_attr_init ( &pthread_attr );
assert ( rc == 0 );
rc = pthread_attr_setstack ( &pthread_attr, targetStack, stackSize );

```



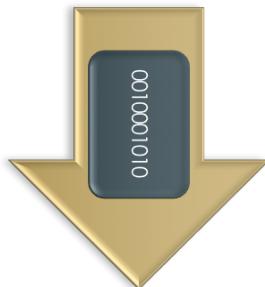
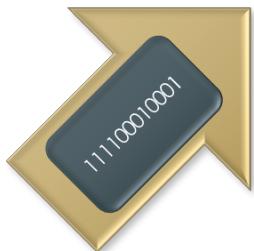
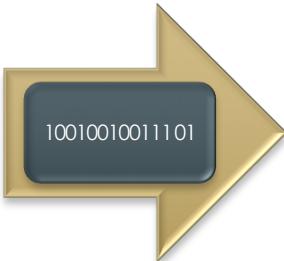
63	67	69	62	73	6f	6e	2f	65	65	63	67	2f	65	63
35	36	38	2f	66	6f	2f	00	73	74	61	63	6b	53	63
74	75	70	2e	63	00	2f	55	73	65	72	73	2f	63	67
52	73	6f	6e	2f	65	65	63	67	2f	65	63	65	35	36
66	6f	6f	2f	73	74	61	63	6b	53	65	74	75	70	72
00	5f	6c	61	62	5f	6d	61	69	6e	5f	5f	5f	66	75
64	00	5f	6d	61	69	6e	00	5f	5f	5f	5f	66	75	6e
5f	2e	33	37	30	31	00	5f	5f	5f	5f	5f	66	75	6e
3e	33	37	32	35	00	6d	61	69	6e	2e	63	00	2f	63
72	73	2f	63	67	69	62	73	6f	6e	2f	65	65	61	61
65	63	65	35	36	38	2f	66	6f	6e	2f	6d	6d	5f	5f
6f	00	5f	6c	61	62	5f	6d	61	69	6e	00	5f	70	76
55	6e	63	5f	5f	2e	33	37	32	35	00	5f	4e	58	41
6e	63	5f	5f	2e	33	37	32	35	00	5f	4e	58	41	00
5f	5f	6d	60	5f	4e	58	41	72	67	63	00	5f	68	65
64	65	72	68	5f	5f	5f	70	72	6f	6e	61	6d	65	61
62	5f	6d	60	5f	65	78	65	63	75	74	65	5f	68	65
5f	74	6d	61	69	6e	76	69	72	6f	6e	00	5f	72	74
61	72	74	00	5f	61	64	00	5f	6c	61	62	5f	6d	6d
6e	00	5f	65	78	69	74	00	5f	67	65	72	74	00	5f
73	69	7a	65	00	5f	61	73	73	65	72	74	70	74	00
61	70	00	5f	70	72	69	6e	74	66	00	5f	70	74	00
65	61	64	5f	61	74	74	72	5f	69	6e	69	74	00	5f
74	68	72	65	61	64	5f	61	74	74	72	5f	73	65	74

```
stackSize = STACK_SIZE;
targetStack = (void *)STACK_LOCATION;
pagesize = getpagesize();
mmap_result = NULL;
lab_main_thread_id;
pthread_attr;
rc = 0;
args = { argc , argv };
lab_main_return = NULL;

// check that our target stack size is an integer multiple of pagesize
stackSize += pagesize - ( stackSize % pagesize );

// Allocate a new stack at a fixed location
mmap_result = mmap(targetStack, ( stackSize + pagesize ),
PROT_READ|PROT_WRITE|PROT_EXEC, MAP_ANON|MAP_FIXED);
assert ( mmap_result != MAP_FAILED );

// Create a new pthread to run lab_main()
rc = pthread_attr_init ( &pthread_attr );
assert ( rc == 0 );
rc = pthread_attr_setstack ( &pthread_attr, targetStack, stackSize );
assert ( rc == 0 );
```



Data Representation

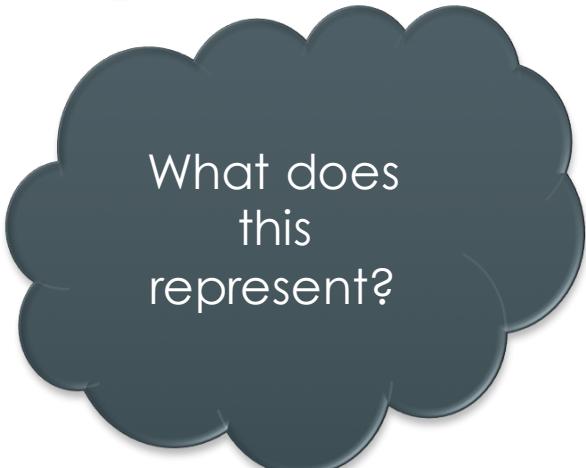
Because everything (instructions, input and output) is represented as numbers, we need to be careful.

Try to avoid thinking of memory only in terms of how you define your variables:

```
int x = 1145258561;
```

01000001010000100100001101000100

Base 2



What does
this
represent?

0100000101000100100001101000100

Base 2



0100 0001 0100 0010 0100 0011 0100 0100

Nibbles



0x41424344

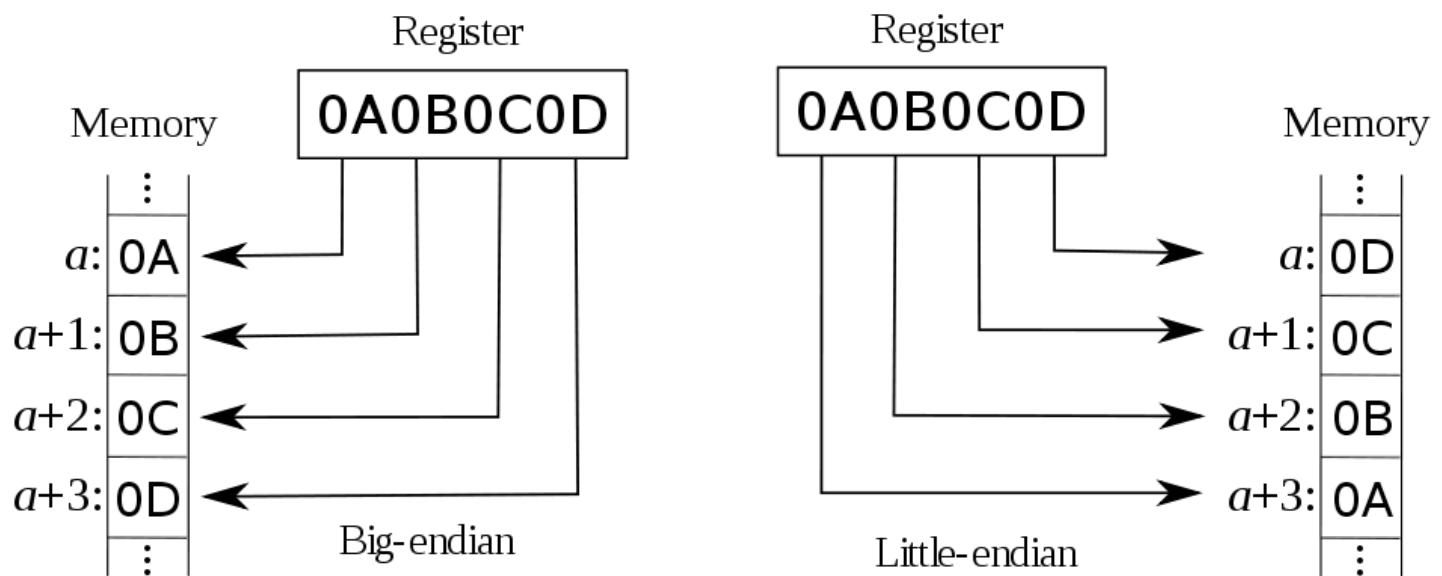
Base 16



1094861636

Base 10

Big-Endian vs. Little-Endian



Intel is Little-Endian

0100000101000100100001101000100

Base 2



0100	0001	0100	0010	0100	0011	0100	0100
------	------	------	------	------	------	------	------

Big-Endian



0x41424344

Little-Endian



0x44434241

Base 16



1094861636

Base 10

Base 16



1145258561

Base 10

0100000101000100100001101000100

Base 2



0

1000001

010000100100001101000100

IEEE-754



781.035217

Base 10

0100000101000100100001101000100

Base 2



01000001

01000010

01000011

01000100

Bytes

65

66

67

68

Base 10



'A'

'B'

'C'

'D'

ASCII

0100000101000100100001101000100

Base 2



01000001

01000010

01000011

01000100

Bytes

65

66

67

68

Base 10



65.66.67.68 → swbell.net

IPv4

0100000101000100100001101000100

Base 2



01000001

01000010

01000011

01000100

Bytes

65

66

67

8-bit Color



Pixel

0100000101000100100001101000100

Base 2



01000001

01000010

01000011

01000100

Bytes



01000001

rex.B

01000010

rex.X

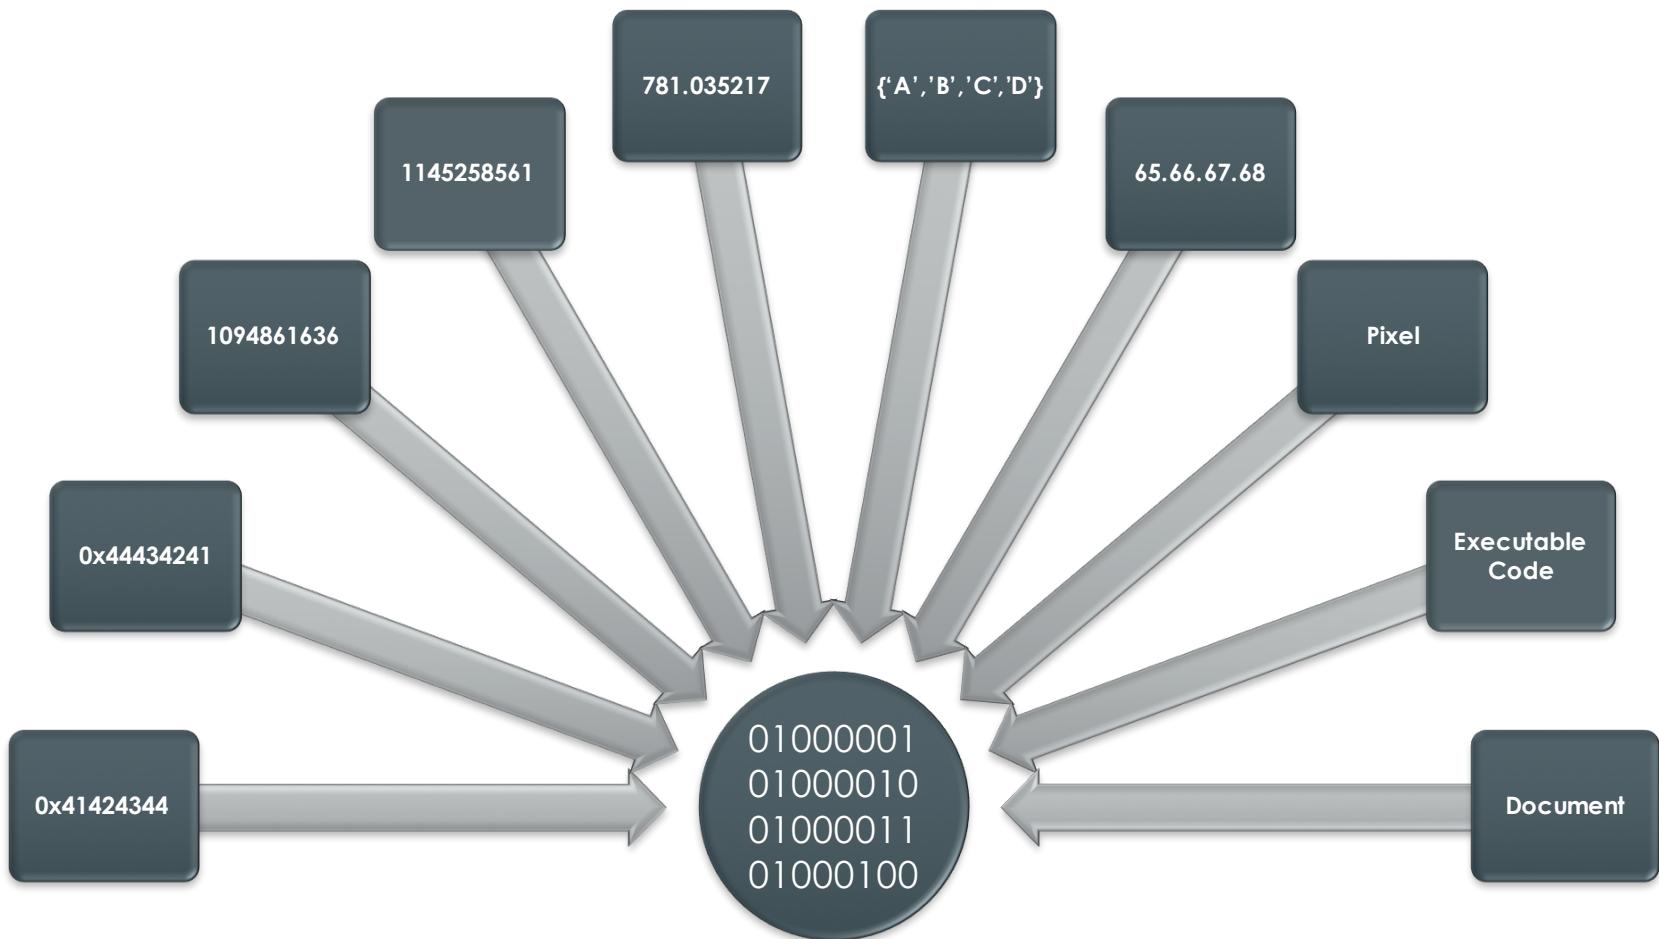
01000011

rex.XB

01000100

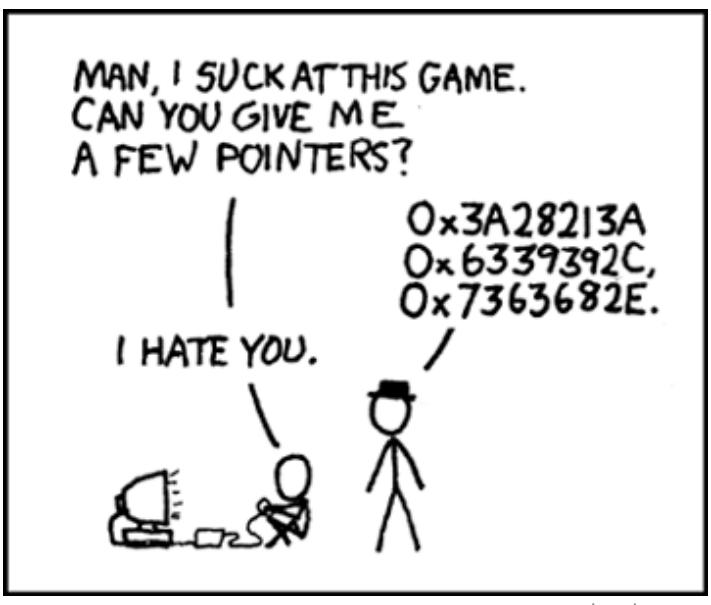
rex.R pop

Intel Assembly Code

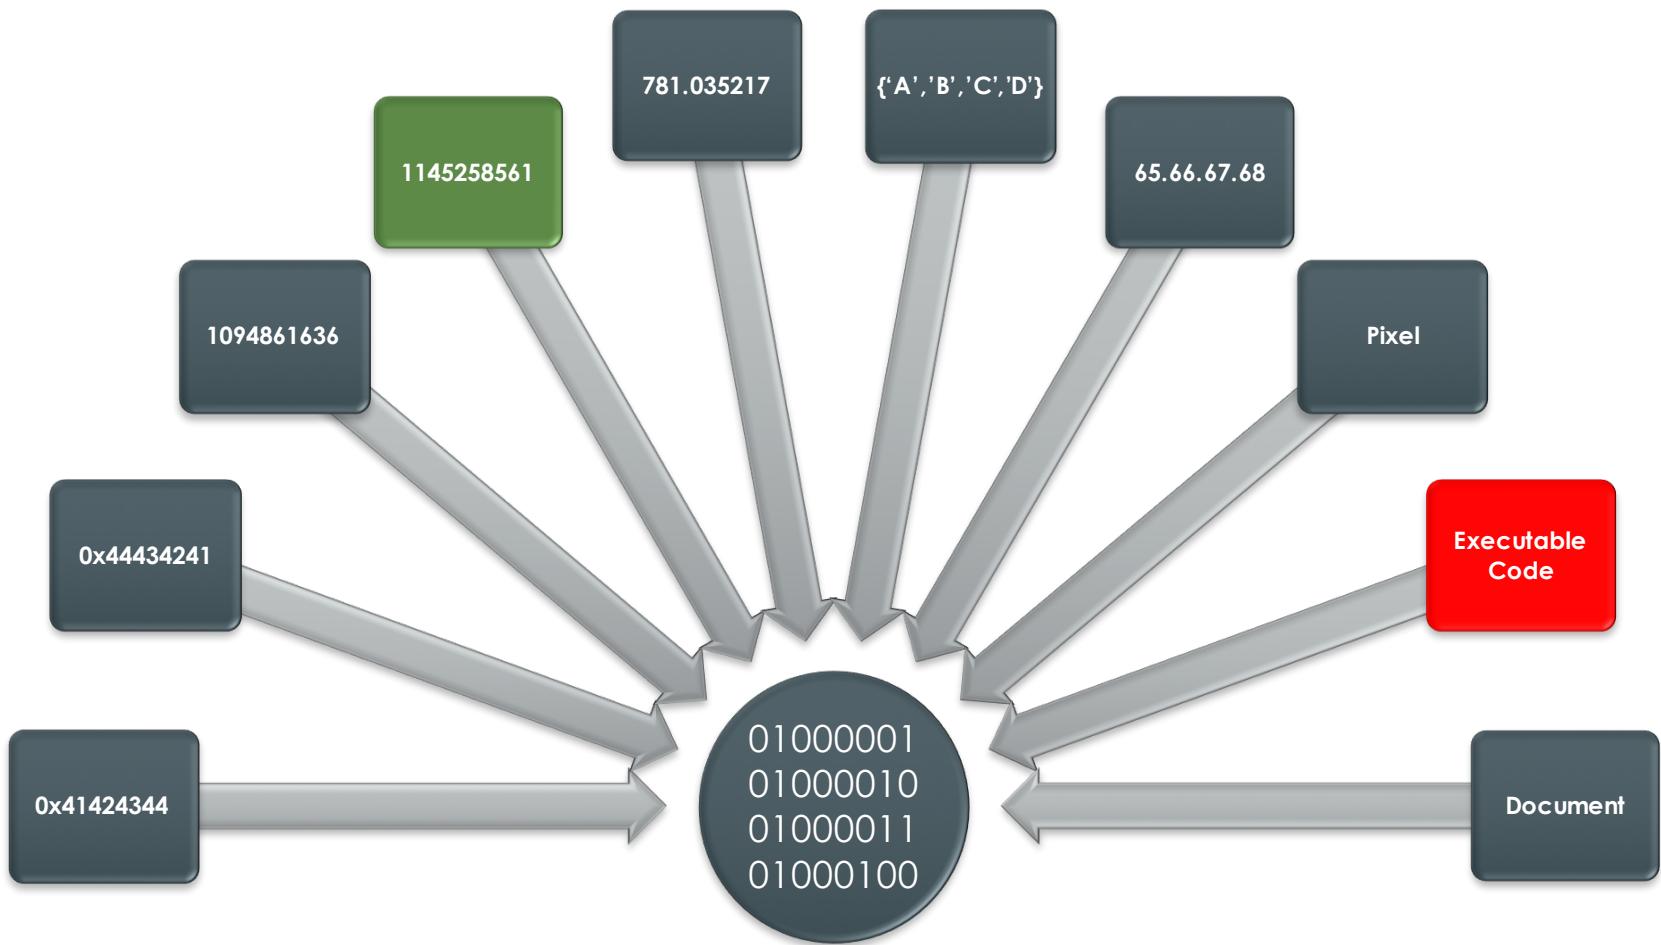


Data Representation

What if our assumptions about how the data will be interpreted are wrong?



xkcd.com

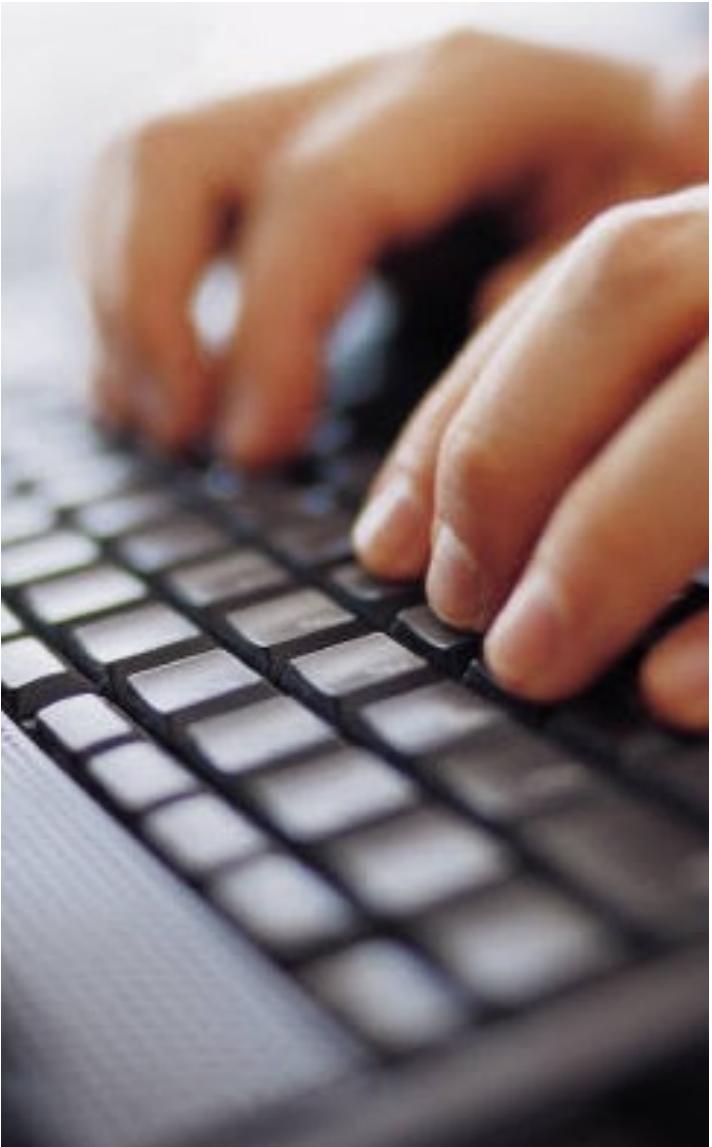


**Many security exploits rely on data being interpreted
in a different way than was originally intended.**

```
unsigned long int i_hello = 0x6f6c6c6548;  
unsigned long int i_world = 431316168567;  
  
printf("%s %s", (char *)&i_hello, (char *)&i_world);
```

ASCII	Hex	Dec	ASCII	Hex	Dec
'A'	41	65	'a'	61	97
'B'	42	66	'b'	62	98
'C'	43	67	'c'	63	99
'D'	44	68	'd'	64	100

Remember: Intel is Little-Endian; the bytes in an integer are stored in reverse order in memory.



Assembly Code

Intel Assembler, gdb

Intel Assembler

We will be making use of Intel assembler, both to discuss and to craft various security exploits. You will not need to be an expert, but should be comfortable in reading assembly code.

For reference, see:

<http://www.cs.virginia.edu/~evans/cs216/guides/x86.html>

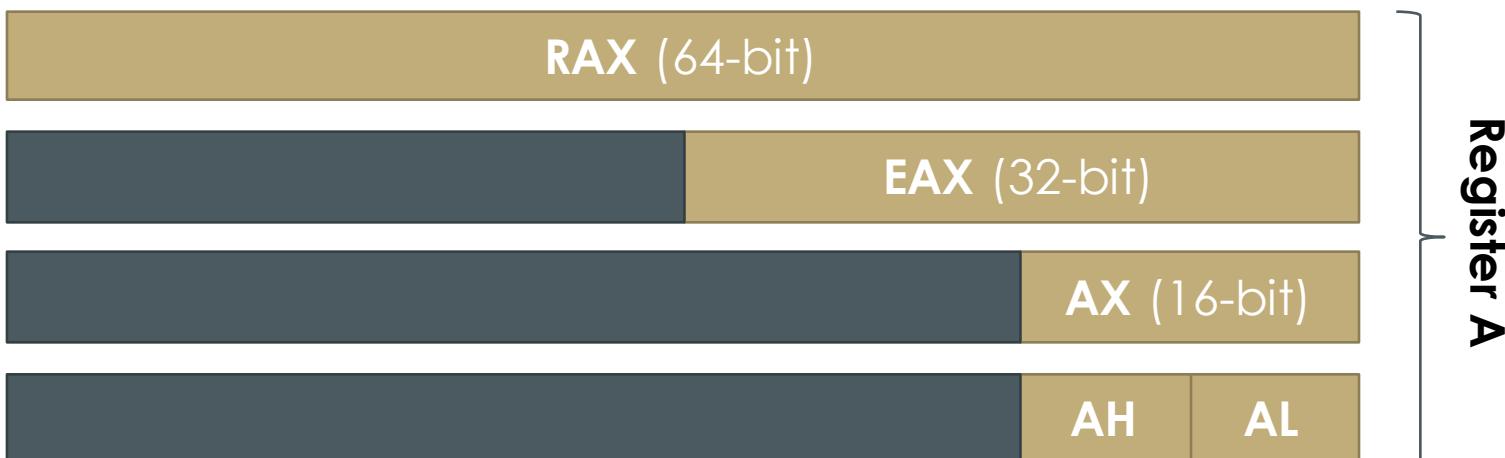
Intel Assembler

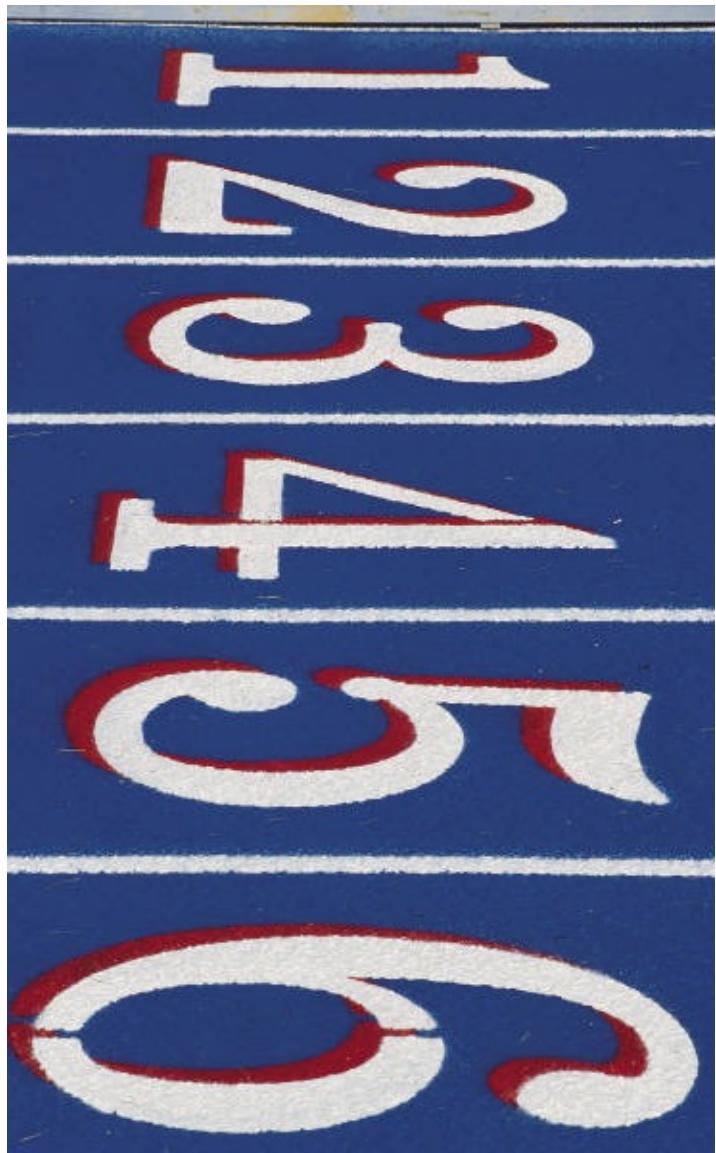
Intel assembly code looks very similar to most other assembly languages you may already be familiar with (e.g., Motorola 68k)

```
    movl -4(%rbp), %eax  
    addl $1, %eax  
    movl %eax, -16(%rbp)
```

Intel Registers

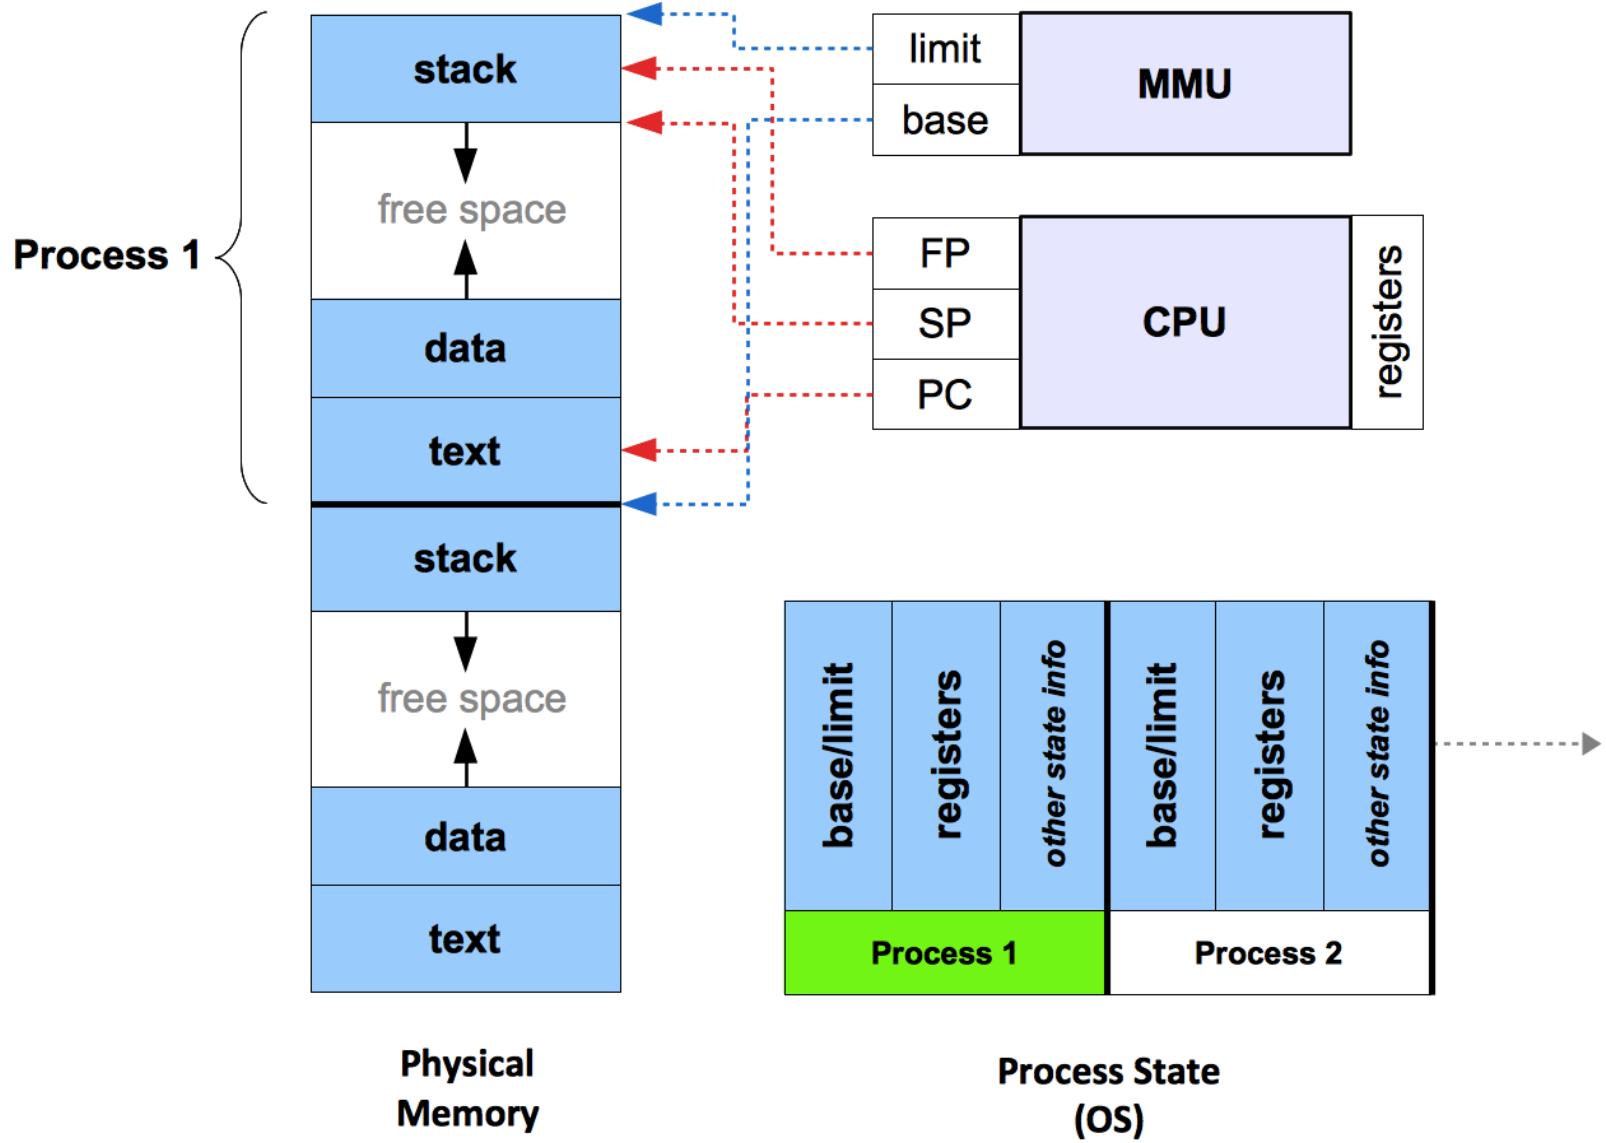
Current Intel CPUs have six general-purpose registers; the various bits in each register can be accessed in a number of ways:



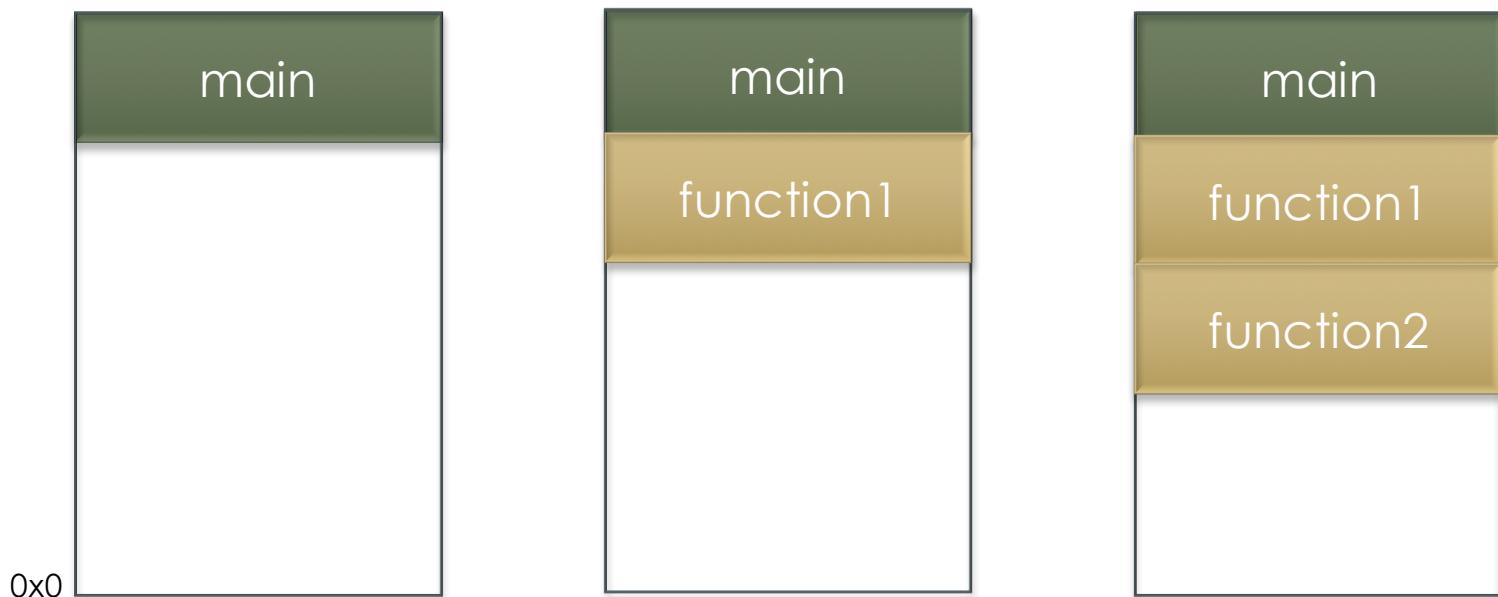


Stacks

Stack Frames, Function Calls



Process Stack Growth



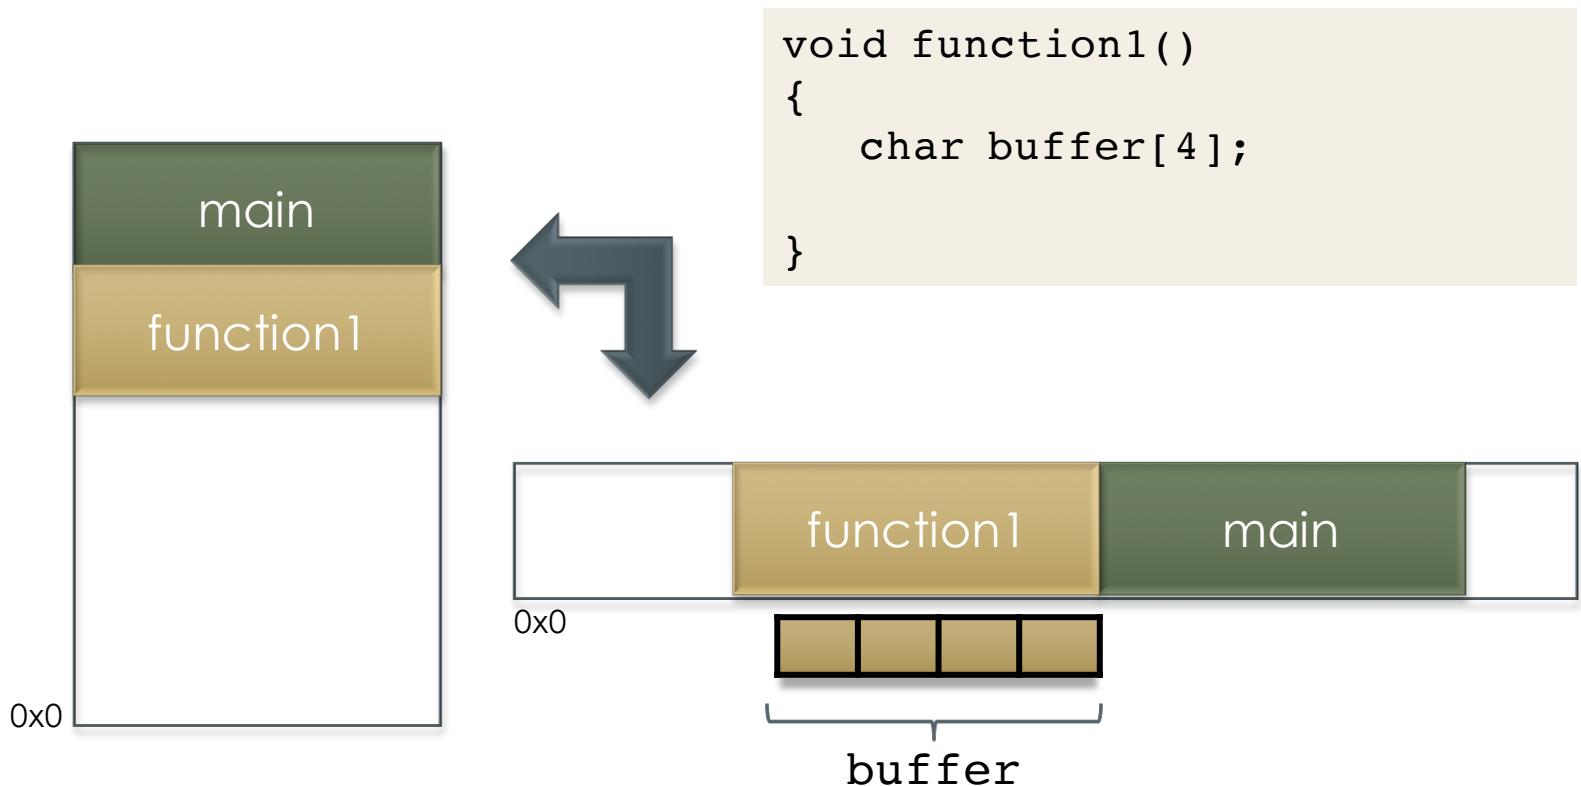
The process stack grows downwards on Intel.

Stack Frames

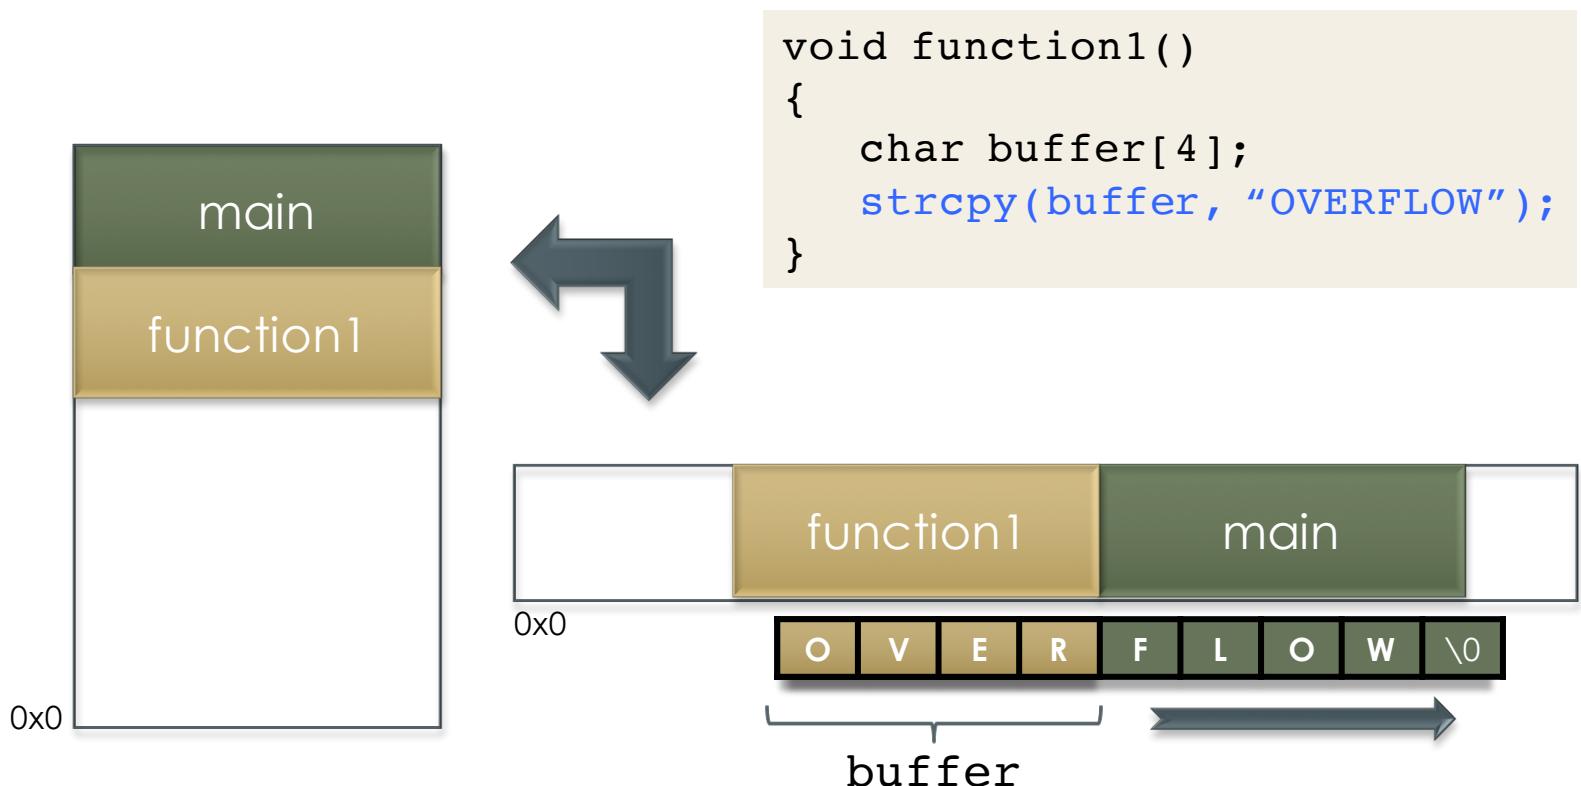
Recall: What information does a stack frame contain?



Process Stack Growth



Process Stack Growth



Using gdb to Examine Stacks

break

Defines a new breakpoint

run

Starts a new process

where

List of current stack frames

up / down

Move between frames

info frame

Display info on current frame

info args

List function arguments ("gcc -g")

info locals

List all local variables ("gcc -g")

print

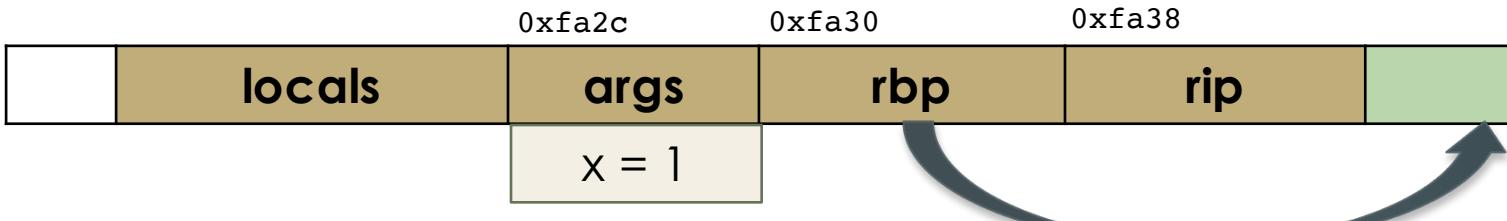
Display a variable

x

Display the contents of memory

gdb: info frame

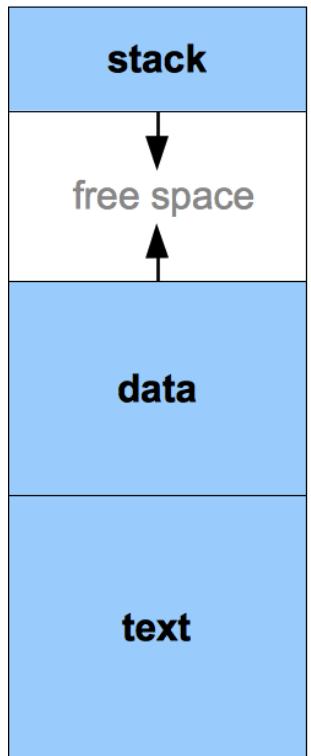
```
(gdb) info frame
Stack level 0, frame at 0xfa40:
rip = 0x100000ec7 in foo (foo.c:3); saved rip 0x100000ef9
called by frame at 0xfa70
source language c.
Arglist at 0xfa38, args: x=1
Locals at 0xfa38, Previous frame's sp is 0xfa40
Saved registers:
    rdi at 0xfa2c, rbp at 0xfa30, rip at 0xfa38
(gdb)
```



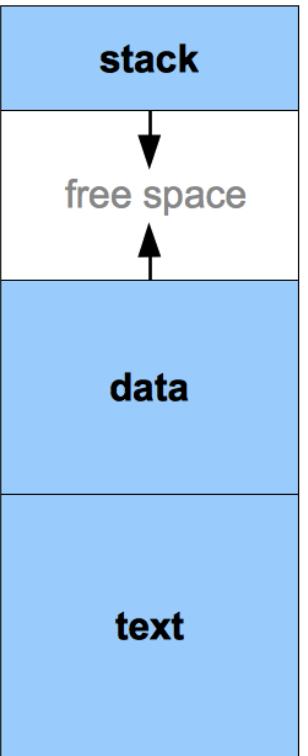


Processes

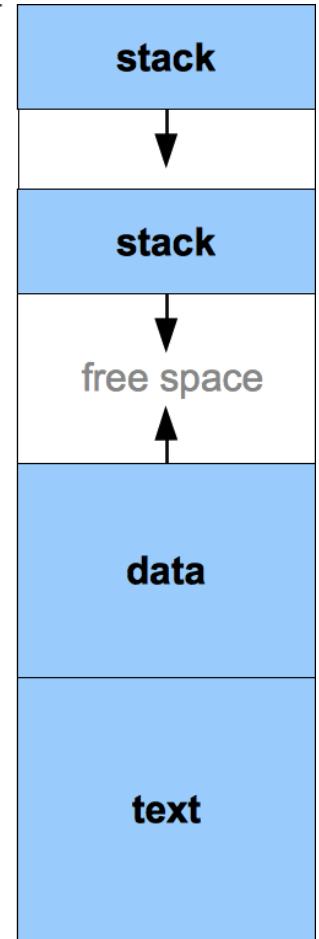
Threads, fork, exec



2 Processes



Process Address Space



2 Threads

